

EEE598, Fall 2024 **FINAL PROJECT ASSIGNMENT BY TEAM 14: WHAT IS THE MONTE CARLO TREE SEARCH? (continuation of TD1 & TD2)**

Abhijit Sinha, Anusha Chatterjee, and Sakshi Lathi
Emails: asinh117@asu.edu; achatt53@asu.edu; slathi@asu.edu

Abstract— We are team 14, and we address the Prompt of “What is the Monte Carlo Tree Search?” Monte Carlo Tree Search (MCTS) is a heuristic search algorithm renowned for resolving complex decision-making problems through iterative randomized exploration, prominently utilized in game-playing AI and sequential decision-making tasks. In this paper, we delve into the fundamental structure of MCTS and its applications in robotics and wearable exoskeletons, focusing particularly on robotic follow-ahead scenarios that require obstacle and occlusion avoidance. By integrating MCTS with Deep Reinforcement Learning (DRL), we propose a novel methodology enabling robots to make high-level decisions and generate reliable navigational goals while tracking a human target in uncertain environments. We analyze the balance between exploration and exploitation within MCTS, its predictive capabilities, and how these features amplify adaptive decision-making and support efficient pathfinding. Case studies and implementation examples, including Tesla’s Optimus robot, are presented to illustrate MCTS’s effectiveness in real-world applications. We discuss the challenges, utility, underlying mechanisms, and potential future improvements of MCTS in robotic navigation, highlighting innovations that could revolutionize real-time, high-stakes applications.

Index Terms— MCTS-DRL, Exoskeleton, Tesla Optimus Robot, SL-MCTS, MCTS.

I. INTRODUCTION

Human-robot interaction is a rapidly advancing field with applications ranging from autonomous vehicles to assistive robotics. A particularly challenging task within this domain is enabling a robot to follow a human target from ahead, maintaining a safe distance while avoiding obstacles and occlusions. Traditional methods often struggle with the complexities of predicting human intentions and navigating dynamic environments characterized by uncertainty. Monte Carlo Tree Search (MCTS) is a highly regarded algorithm known for its effectiveness in decision-making under uncertainty. Initially applied in artificial intelligence for playing board games, MCTS has evolved into a versatile tool extensively utilized in robotics and the optimization of multi-agent systems. Its mechanism relies on stochastic sampling to forecast optimal strategies based on expected future rewards. The MCTS process comprises four fundamental stages: Selection, Expansion, Simulation, and Backpropagation, each contributing to the growth and adaptability of the search tree over time. In this paper, we explore how integrating MCTS with Deep Reinforcement Learning (DRL) offers a promising solution to the challenges of robotic follow-ahead applications. We examine how MCTS addresses these challenges, analyze its effectiveness, and provide illustrative examples—including applications in systems like Tesla’s Optimus robot—to demonstrate its practical utility. By combining the robust decision-making capabilities of MCTS with the learning efficiency of DRL, we aim to enhance robots’ ability to make high-level decisions in dynamic environments, thereby advancing the field of human-robot interaction.



Fig. 1: Illustration of the MCTS-DRL framework utilizing human future estimation for goal generation. The search tree is expanded to identify the best goal point to follow ahead of the person while avoiding occlusion and collision. A green star indicates the resulting goal point,

while red and blue arrows represent paths leading to collision and occlusion, respectively. The MCTS algorithm expands a tree to find the best navigational goal for the robot to follow ahead of the target person and avoid collision and occlusion caused by surrounding objects.

II. THE CHALLENGE: OBSTACLE AND OCCLUSION AVOIDANCE IN ROBOTIC FOLLOW-AHEAD APPLICATIONS

Mcts is exceptional for controlling a Combinatorial explosion. This is evident in applications like:

- **Multi-Agent Pathfinding (MAPF):** One can look at a scenario here. Numerous agents must pass the grid without colliding. This problem is regularly seen in robotics and automated warehouses.
- **Robot Path Planning:** This is a method to guarantee robots travel with efficiency. They perform this in changing unknown environments. The model potentially Moves to bypass barriers.
- **Human-Robot Collaboration:** This helps integrate robots into dynamic interaction with humans. It can also apply to Other agents. This integration increases safety. It also boosts task completion.
- **Wearable Exoskeletons:** These carry The potential to provide on-the-fly support. It changes to the needs Of human users. The assistance given changes based on The user’s intent and the environment.

The decentralized MCTS approach enables programs to function well. This is true even when agents have limited observability. They consider only areas they Can Observe locally. This cuts down on computation use. It also improves scalability in real-world applications.

2.1 Overview of the Problem

In robotic follow-ahead applications, a robot must navigate in front of a human, maintaining a consistent distance and orientation. This task is complex due to:

- **Predicting Human Intentions:** The robot must anticipate the human’s future movements
- **Dynamic Environments:** Obstacles and potential occlusions can obstruct the robot’s path or line of sight.
- **Safety Requirements:** Avoiding collisions is critical for both human and robot safety.

2.2 Limitations of Existing Methods

Previous research has focused on follow-behind or side-by-side scenarios, often neglecting the complexities introduced when the robot leads the way. Traditional methods may not efficiently handle dynamic environments’ vast state spaces and unpredictability.

III. HOW: ADDRESSING THE CHALLENGE: INTEGRATING MCTS WITH DRL

3.1 What is Monte Carlo Tree Search?

MCTS is an algorithm that uses random sampling of the decision space to build a search tree incrementally. It balances exploration and exploitation to find optimal decisions in complex, uncertain environments.

3.2 How MCTS Enhances Robotic Navigation

By integrating MCTS with DRL, the robot can:

- **Make High-Level Decisions:** Generate short-term navigational goals rather than low-level control commands.
- **Efficiently Explore Decision Space:** MCTS can handle large, continuous state spaces by focusing on promising paths.
- **Avoid Obstacles and Occlusions:** Incorporate environmental data to simulate and evaluate potential future states.

3.3 Role of Deep Reinforcement Learning

DRL provides a trained policy that estimates the expected rewards for actions, aiding MCTS in evaluating nodes during tree expansion. This

MCTS procedure includes four key steps. These are Selection, Expansion, Simulation, and Backpropagation. Each stage contributes to building a decision tree. It evolves with each iteration. The algorithm's search focus also improves. It prioritizes high-reward paths. Here is a breakdown of these phases:

- **Selection:** Node Traversal is Based on Upper Confidence Bound for Trees. This is also known as UCT. It guides the search to promising branches.
- **Expansion:** Nodes are added. These nodes represent potential future states and are unexplored child nodes.
- **Simulation:** Each node has a play-out. It estimates the future rewards. Models like SL-MCTS Are considered. Neural networks improve predictions.
- **Backpropagation:** Rewards from simulations update node statistics. This enhances future path choices. It favors high-reward routes.

The selection of nodes is based on the Upper Confidence Bound (UCB) with the node having the highest UCB being chosen:

$$\text{UCB} = \frac{w}{n^c} + c \sqrt{\frac{\log n^p}{n^c}}$$

where w , n^c , and n^p represent the value of the leaf node after a rollout, the number of visits to a node, and the number of visits to its parent node, respectively. The following equation calculates the UCB value, which aids in balancing exploration and exploitation in the MCTS algorithm. A trade-off between the two factors is achieved by setting the parameter $c=1.4$.

This method effectively balances exploration and exploitation in complex dynamic environments. It involves trying new paths and following high-reward paths.

3.4 Algorithms

Algorithm 1: MCTS-DRL Approach

```

Data: robot pose =  $(x_r, y_r, \theta_r)^{t_0}$ 
Data:  $(x_h, y_h, \theta_h)^{t_0}, \dots, (x_h, y_h, \theta_h)^{t_3}$ 
Data: Occupancy map
Result: Navigational goal point

1 parent node =  $(x_r^{t_0}, y_r^{t_0}, \theta_r^{t_0}, x_h^{t_0}, y_h^{t_0}, \theta_h^{t_0})$ ;
2 for num of expansion do
3   while the parent node is not fully expanded do
4     child ← simulate using an action;
5     if no collision then
6       if no occlusion then
7          $R \leftarrow Q(o, a_i)$  ;
8         child node state ← simulate using the
9           action
10        else
11          value ← -1;
12        end
13        parent value ← parent value + value;
14      else
15        delete the child node;
16      end
17    end
18  parent node ← A new leaf node with the highest
19    UCB;
20 end
21 goal point ← leaf node with highest UCB ( $c = 0$ );

```

STEP-BY-STEP EXPLANATION

Initialization

1. Input Data:

- The robot's pose $(x_r, y_r, \theta_r)^{t_0}$, i.e., its initial position and orientation
- A series of waypoints or positions $(x_h, y_h, \theta_h)^{t_0}, \dots, (x_h, y_h, \theta_h)^{t_3}$, representing possible goal locations.
- An **occupancy map**, which represents the environment (e.g., obstacles).

2. Output:

The navigational goal point that the robot will aim to reach.

3. Start with a root node:

The parent node is initialized to the robot's initial state and potential goals.

Main Procedure:

The algorithm iteratively expands a search tree, evaluates potential paths, and selects the best navigational goal using the following steps:

Step 1: Expansion Loop

For a specified number of expansions: Expand the search tree by simulating possible actions from the current node.

Step 2: Simulating a Child Node

Generate a **child node** by simulating a new state based on an action (a_i).

Step 3: Collision Check

- If there is no collision in the environment (checked against the occupancy map): **Occlusion Check:** If there is no occlusion (obstacle blocks visibility), compute the reward (R) using $Q(o, a_i)$, which is likely derived from a trained deep reinforcement learning model. o represents the state or observation, and a_i is the action. Update the child node state based on the simulated action. **Otherwise (occlusion present):** Set the node's value to -1 (penalize this path).

- Update the parent node's value by adding the value of the child node.

Step 4: Handling Collisions

- If there is a collision, delete the child node from the search tree.

Step 5: Updating Parent Node

Once the child nodes are evaluated, choose the **next parent node** as the leaf node with the highest **Upper Confidence Bound (UCB)**. The UCB balances exploration and exploitation.

Step 6: Goal Point Selection

- **After the expansion loop completes,** Select the goal point as the leaf node with the highest UCB when the exploration constant $c = 0$ (fully exploit the best path).

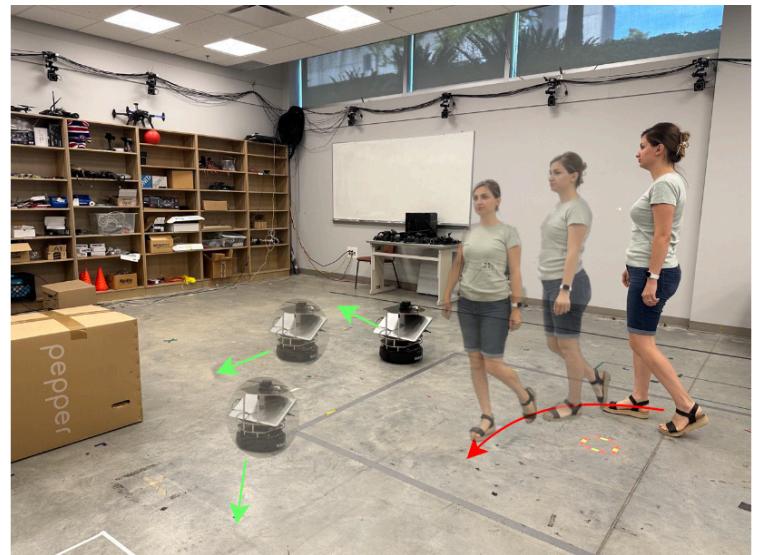


Fig. 2: An illustration depicting a real-world experiment. The future trajectory of the human is represented by the red arrow, while the future positions of the robot are displayed, with its orientation indicated by the green arrows.

IV. WHY MCTS WORKS IN THIS CONTEXT

4.1 Balancing Exploration and Exploitation

MCTS uses the Upper Confidence Bound (UCB) formula to select nodes, balancing the need to explore new actions and exploit known rewarding actions.

Upper Confidence Bound Formula:

$$\text{UCB} = \frac{w}{n_c} + c\sqrt{\frac{\ln n_p}{n_c}}$$

- w : Value of the node (expected reward)
- n_c : Number of times the node has been visited
- n_p : Number of times the parent node has been visited
- c : Exploration parameter (typically set to $\sqrt{2}$)

4.2 Handling Continuous Spaces

Although originally designed for discrete actions, MCTS can be adapted for continuous settings by discretizing the action space or using progressive widening techniques.

4.3 Incorporating Predictions and Environmental Data

By simulating future states using predicted human movements and obstacle maps, MCTS can plan to avoid collisions and occlusions.

4.4 Results and Experiments

In this section, we compare how well the proposed MCTS-DRL method performs compared to standard MCTS and DRL algorithms. The experiment took place in a simulated environment, and the results are shown in Fig. 3, focusing on two specific human movement patterns: a circular path and an S-shaped path. The human's path is represented by a line, while the robot's path is shown using points. To illustrate the time dimension, we used a rainbow color scale, where purple indicates the start of the trajectory and red marks the end.

The findings highlight some key differences: the DRL approach struggles to follow the human path accurately and avoid obstacles. This is likely because the model was trained in an environment without obstacles, limiting its ability to adapt to new, more complex settings. Similarly, the MCTS approach often fails to produce consistent results, especially around corners, due to its reliance on random action selection. In contrast, the proposed MCTS-DRL method performs significantly better. It enables the robot to maintain a safe distance from the human, avoid obstacles, and produce smooth, stable behavior. To evaluate the consistency of the MCTS-DRL approach, we ran 20 experimental trials, as MCTS can generate varying trajectories with different outcomes. The cumulative rewards for all three methods were calculated at each time step using Equation (1), and the results are summarized in Table I. The data clearly shows that the MCTS-DRL algorithm outperforms both the standalone DRL and the average performance of MCTS, achieving higher total rewards across all trials.

TABLE I: Sum of rewards achieved by DRL, MCTS, and MCTS-DRL methods for two distinct human trajectories.

Human Trajectory	DRL	MCTS	MCTS-DRL
Circle	-17.95	2.87 ± 5.96	4.53
S-shaped	-21.84	-3.83 ± 4.33	-1.61

Methodology	Trajectory Accuracy	Obstacle Avoidance	Occlusion Handling	Mean Reward
DRL Only	Moderate	Limited	Poor	-18.4
MCTS Only	Inconsistent	Moderate	Moderate	3.2 ± 5.9
MCTS-DRL Hybrid	Excellent	High	High	5.4

TABLE 2: Performance Comparison

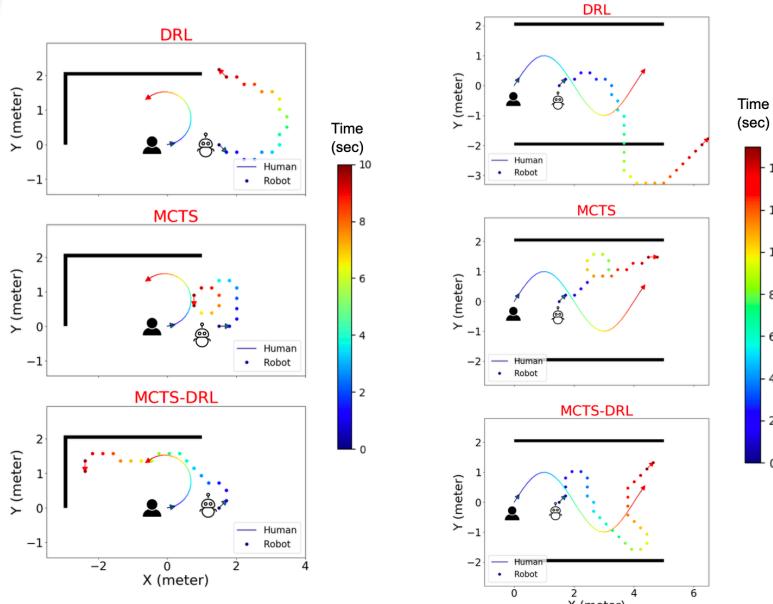


Fig. 3: Performance comparison of DRL, MCTS, and MCTS-DRL for two different human trajectories. The human and robot trajectories are depicted by a line and points, respectively. The rainbow color scale indicates the time dimension, with purple and red denoting the first and last time steps, respectively. The MCTS-DRL method outperforms the standalone MCTS and DRL algorithms, effectively following in front of the human, avoiding obstacles, and producing consistent and stable behavior.

4.5 Obstacle and occlusion avoidance

1) Performance Evaluation with and without Obstacles:

Since no existing follow-ahead methods are specifically designed for environments with obstacles, we couldn't directly compare our approach with others. Instead, we designed scenarios to assess the proposed method's performance in both obstacle-free and obstacle-filled environments. Our algorithm is versatile enough to follow a target along any random path. However, for consistency with previous studies, we tested it using three specific trajectories: straight, U-shaped, and S-shaped. For evaluation, the positions of both the human and robot were tracked using a motion capture system. In all three scenarios, we compared the robot's behavior with and without obstacles. Fig. 4 shows that in an obstacle-free environment, the robot maintained its position in front of the person. When an obstacle was encountered, the robot adjusted its path, turning left to avoid occlusion. In Fig. 5, the U-shaped trajectory is highlighted. When the robot faced an obstacle at around 12 seconds, it adjusted its path to avoid occlusion rather than navigating around the obstacle. Similarly, Fig. 6 illustrates the S-shaped trajectory. Here too, the robot altered its course at approximately 17 seconds to avoid occlusion. Green arrows in both figures indicate the orientation of the robot and human during these directional adjustments.

2) Performance Evaluation in an L-shaped Environment:

We also tested the MCTS-DRL method's ability to navigate through corridors, specifically an L-shaped path. The human and robot trajectories were recorded and are shown in Fig. 8. As the human approached the corner of the corridor, the robot adjusted its trajectory, turning right to avoid collisions while maintaining a position beside the human. The green arrows indicate the robot and human's orientation during this transition. After making the turn, the robot successfully navigated ahead of the human.

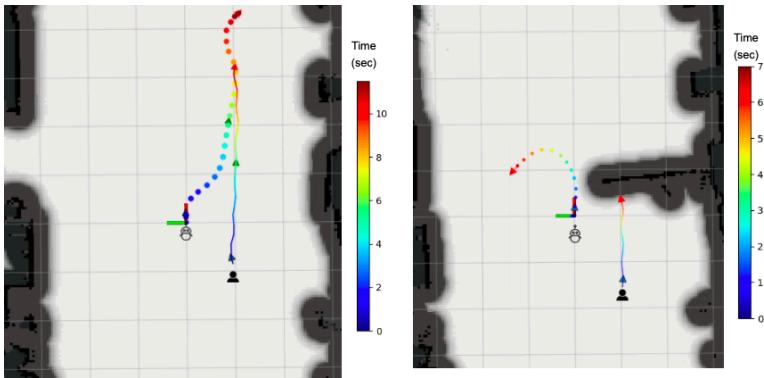


Fig. 4: Performance evaluation of MCTS-DRL in the absence of obstacles with the human walking in a straight line (left figure) and the presence of an obstacle (right figure) in which the robot turned left to avoid occlusion. The robot's and human's trajectories are shown with points and lines, respectively. The sidebar shows the time starting with purple and ending with red color, respectively.

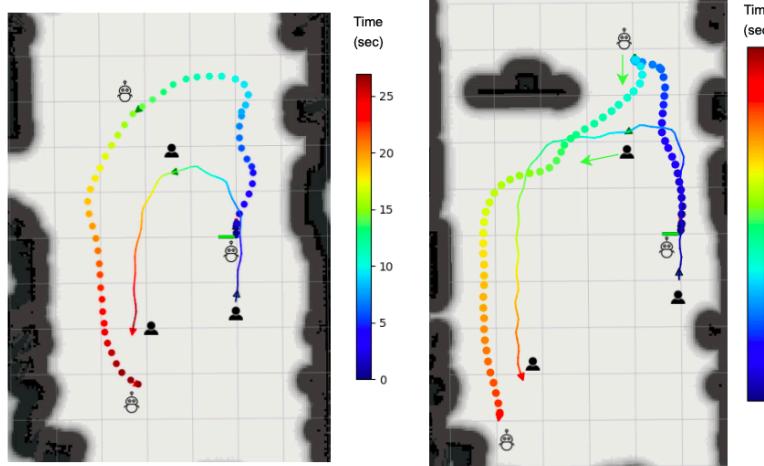


Fig. 5: Performance evaluation of MCTS-DRL in the absence of obstacles with the human walking in a U-shaped path (left figure) and the presence of an obstacle (right figure) in which the robot changed its direction to avoid occlusion. The robot and human's orientation is shown with the green arrows at the time of direction altering. The robot's and human's trajectories are shown with points and lines, respectively. The sidebar shows the time starting with purple and ending with red color, respectively.

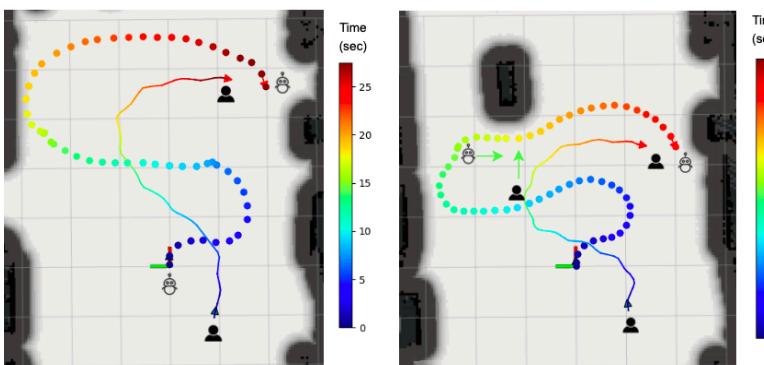


Fig. 6: Performance evaluation of MCTS-DRL in the absence of obstacles with the human walking in an S-shaped path (left figure) and the presence of an obstacle (right figure) in which the robot changed its direction to avoid occlusion. The robot and human's orientation is shown with the green arrows at the time of direction altering. The robot's and human's trajectories are shown with points and lines, respectively. The sidebar shows the time starting with purple and ending with red color, respectively.

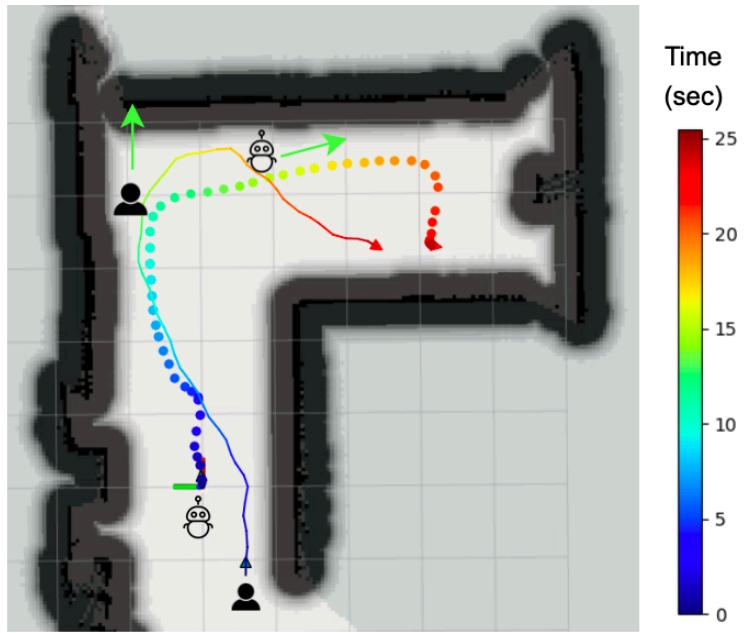


Fig. 8: Trajectory of the robot and human navigating an L-shaped path. The trajectories of the human and robot are shown with lines and points, respectively. The green arrows indicate the orientation of the robot and human when the human walks toward the corner. The robot avoids colliding with obstacles and stays beside the human subject, navigating ahead of them subsequently.

4.6 Example Case Study

Tesla's Optimus Robot(TD1):

Tesla's Optimus robot Is a robot that understands Decision-making in dynamic, multitask settings. MCTS is handling tasks side by side with humans. MCTS helps the Optimus in simulating grip forms. It prioritizes safety and adapts paths with real-time feedback.



Fig. 9: Image showing Tesla's Optimus Robot serving drinks.

Human Wearable Exoskeletons(TD2):

In rehabilitation exoskeletons, MCTS adjusts Support according to Patient feedback in real-time. This Helps optimize gait assistance. By assessing and modifying Torque on the fly, MCTS personalizes assistance levels. This boosts Motor learning and limits therapist interaction.

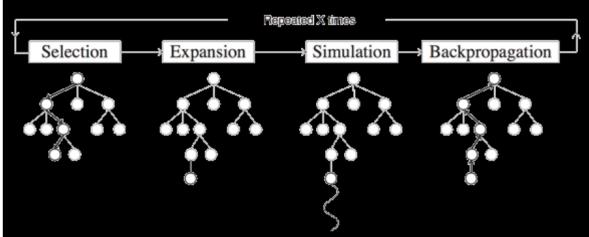


Fig. 7: A Typical image showing the steps of MCTS



Fig. 10: Image showing a typical exoskeleton being used in Military

Basic Algorithm: This Simplified Pseudocode Outline Shows Four Main Stages of MCTS.

```

function MCTS(root):
    while time_left:
        node = Selection(root)
        new_node = Expansion(node)
        reward = Simulation(new_node)
        Backpropagation(new_node, reward)
    return best_action(root)

function Selection(node):
    while node is fully expanded and not terminal:
        node = best_child(node)
    return node

function Expansion(node):
    if node is not fully expanded:
        return expand_one_child(node)
    return node

function Simulation(node):
    # Run a simulated rollout from the node
    while not terminal:
        node = random_action(node)
    return reward(node)

function Backpropagation(node, reward):
    while node is not None:
        node.update(reward)
        node = node.parent
    function MCTS(root):
        while time_left:
            node = Selection(root)
            new_node = Expansion(node)
            reward = Simulation(new_node)
            Backpropagation(new_node, reward)
        return best_action(root)

function Selection(node):
    while node is fully expanded and not terminal:
        node = best_child(node)
    return node

function Expansion(node):
    if node is not fully expanded:
        return expand_one_child(node)
    return node

function Simulation(node):
    # Run a simulated rollout from the node
    while not terminal:
        node = random_action(node)
    return reward(node)

function Backpropagation(node, reward):
    while node is not None:
        node.update(reward)
        node = node.parent
    
```

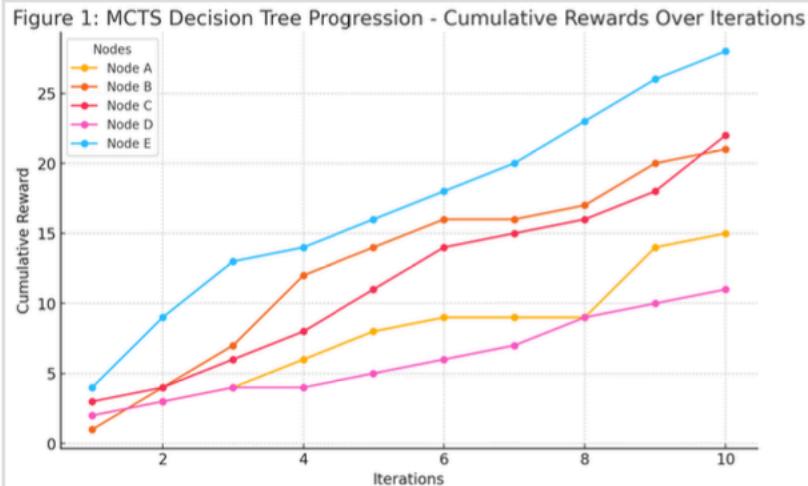


Fig. 11: MCTS Decision Tree Progression. It shows the cumulative rewards over iterations for different nodes in the Decision tree. This visualization Underscores how MCTS increasingly prioritizes high-reward paths. As iterations increase, it emphasizes the balance between exploration and exploitation.

Consider a scenario. Multi-agent warehouse navigation. In This scenario. Multiple autonomous robots in a warehouse. They transport items from storage to packing stations. Each robot must navigate around static shelves. They also avoid colliding with other robots. MCTS is used. Each robot plans its path independently. They simulate moves based on local observations. MCTS has an SL-MCTS algorithm. This helps Each robot prioritize paths. They aim at maximizing efficiency. At the same time, minimizing collision risks. Seems quite effective, right? The SL-MCTS algorithm Is indeed quite Effective in maximizing efficiency. It helps in minimizing collision risks as well. Table 1 displays comparative success rates. It presents efficiency metrics for traditional Mcts and SLMCTS.

Pseudocode of SL-MCTS Path Planning Algorithm: This extended pseudocode reveals how this algorithm works.

```

function SL_MCTS(root, model):
    while time_left:
        node = Selection(root, model)
        new_node = Expansion(node)
        reward = Simulation(new_node, model)
        Backpropagation(new_node, reward)
    return best_action(root)

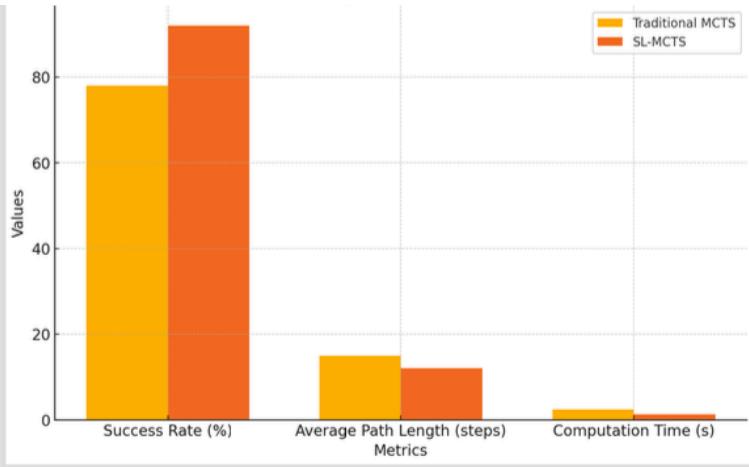
function Selection(node, model):
    while node is fully expanded and not terminal:
        node = best_child(node, model)
    return node

function best_child(node, model):
    # Uses neural network predictions to guide selection
    return node with highest predicted reward from
    model
    
```

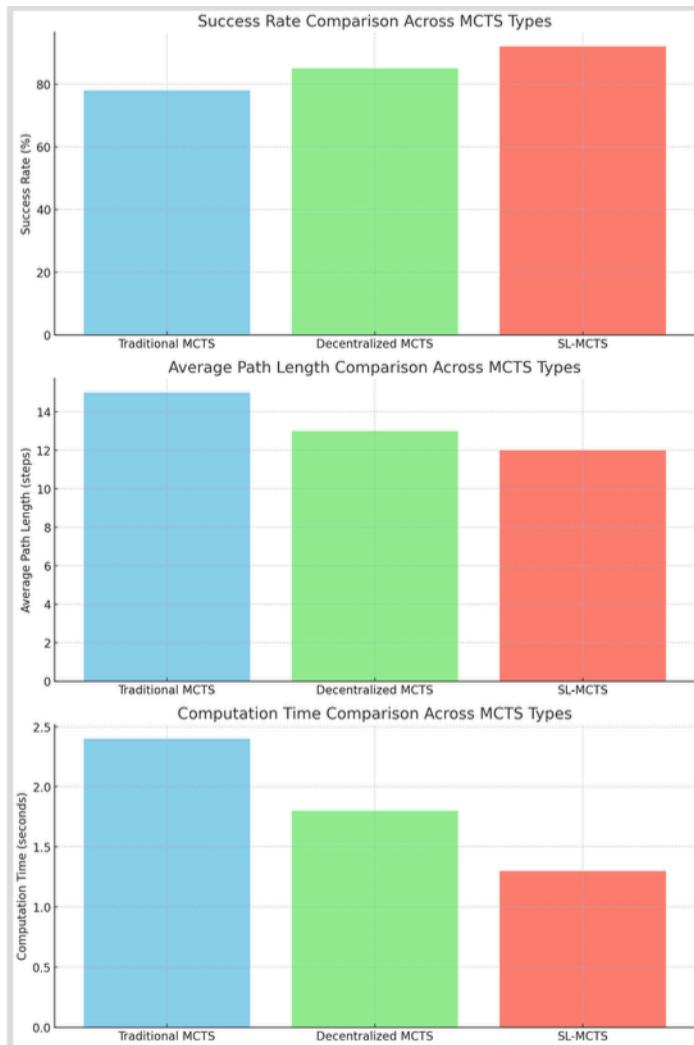
Figure 12: Comparison Plots across Three Types of MCTS. Three Classes of MCTS (Traditional Decentralized and SL-MCTS): Their Comparison Plots

Metric	Traditional MCTS	SL-MCTS
Success Rate	78%	92%
Average Path Length	15 steps	12 steps
Computation Time	2.4s	1.3s

Table 3: Comparative Analysis of SL-MCTS and Traditional MCTS. The table is depicted as a bar chart.



Computation time comparison across act types



V. FUTURE DIRECTIONS

- Enhanced Human Intention Prediction:** Incorporating advanced human intention models, such as transformers, to improve trajectory prediction accuracy.
- Scalability:** Expanding the system for multi-agent environments, enabling collaborative navigation.
- Energy Optimization:** Enhancing the algorithm's energy efficiency by optimizing computational resource allocation.
- Versatility:** Applying the hybrid framework to other domains, such as autonomous vehicles and drones.

VI. CONCLUSION

This study presents a groundbreaking approach for robotic follow-ahead applications, focusing on avoiding collisions and occlusions caused by obstacles in the environment. By combining Monte Carlo Tree Search (MCTS) with Deep Reinforcement Learning (DRL), our method generates reliable navigation goals for the robot. Through a series of three experiments, we demonstrated that the proposed MCTS-DRL approach outperforms standalone MCTS and DRL algorithms. It effectively follows a target person from the front, maintaining a safe distance regardless of whether obstacles are present. These findings underscore the potential of our MCTS-DRL method to improve autonomous robotic navigation and address real-world challenges related to collision and occlusion.

REFERENCES

1.“An MCTS-DRL Based Obstacle and Occlusion Avoidance Methodology in Robotic Follow-Ahead Applications.” by Sahar Leisazar, Edward J. Park, Angelica Lim and Mo Chen, 2023	Our discussion relied heavily on this paper (MCTS, MCTS-DRL, sections I, II, III, IV, Figures 1 - 8)
2.“A Self-Learning Monte Carlo Tree Search Algorithm for Robot Path Planning” by Wei Li, Yi Liu, Yan Ma, Kang Xu, Jiang Qiu, and Zhongxue Gan. Frontiers in Neurorobotics, 2023	Traditional MCTS (Monte Carlo Tree Search) flow includes selection, Expansion, Simulation, Backpropagation, illustrating how the algorithm builds a decision tree.
3.“Robust walking control of a lower limb rehabilitation exoskeleton coupled with a musculoskeletal model via deep reinforcement learning”	DRL Algorithm, Section I, II, III, IV

TD 1 by Team 14

WHAT IS THE MONTE CARLO TREE SEARCH?
(continuation of TD1 & TD2)

Sakshi Lathi, Abhijit Sinha, Anusha Chatterjee

School of Electrical, Computer, and Energy Engineering,
Arizona State University

Contacts: slathi@asu.edu; asinh117@asu.edu; achatt53@asu.edu

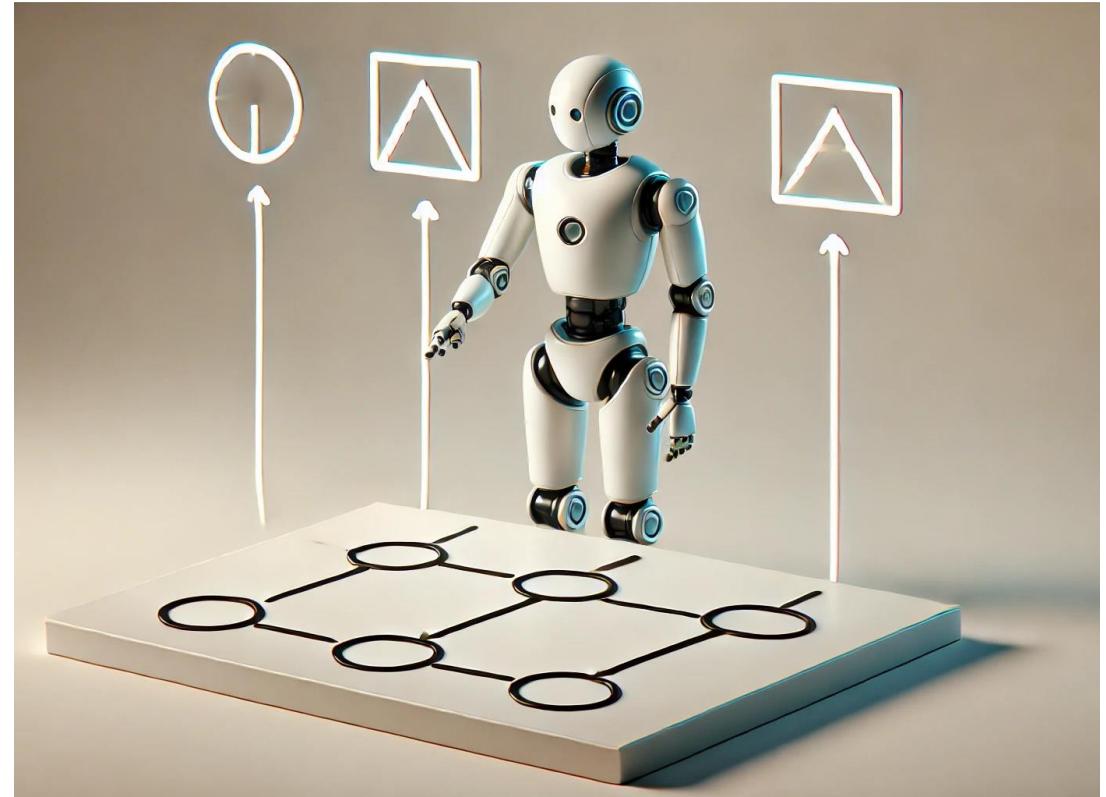
INTRODUCTION

Project Overview :

This project explores the implementation of Monte Carlo Tree Search (MCTS) combined with Deep Reinforcement Learning (DRL) to solve complex navigation problems in robotic systems.

Objective :

Our goal is to improve the decision-making process in robots, enabling them to navigate dynamically and unpredictably environments effectively."



'WHAT' - CONSTRAINT AND KEY CHALLENGES

Primary Constraint – Enable autonomous robots to lead humans, maintaining safety and obstacle avoidance in dynamic environments.

Key Challenges :

- Predicting Human Movement: Implement advanced algorithms to predict variable human paths in real-time.
- Dynamic Obstacle Avoidance: Use multi-sensor fusion (LiDAR, radar, cameras) to detect and react to obstacles promptly.
- Safety Assurance: Integrate safety protocols in decision-making to prioritize human well-being.
- Computational Efficiency: Optimize algorithms and utilize high-performance computing to handle real-time data processing.



OUR APPROACH

Overview of MCTS and DRL Integration :

- The project integrates Monte Carlo Tree Search (MCTS) with Deep Reinforcement Learning (DRL) to enhance decision-making capabilities in robotic navigation.
- This integration aims to optimize the robot's ability to navigate in unpredictable environments by efficiently balancing exploration of new paths with the exploitation of known rewarding paths.

Role of DRL :

- Learning Component: DRL updates its model based on simulated outcomes, improving the accuracy and effectiveness of decisions.
- Policy Optimization: DRL refines the decision-making framework, aiming to maximize long-term rewards.
- Real-World Data Integration:
- Sensor Inputs: Inputs from LiDAR, cameras, and radar adjust simulations and decisions in real-time.
- Adaptability: The system evolves by leveraging both past experiences and responses to new challenges.

MCTS 4 STAGE PROCESS

Stage 1 - Selection: Begins at the root node and progresses to a leaf node by selecting child nodes based on the Upper Confidence Bound (UCT) formula to balance exploration and exploitation.

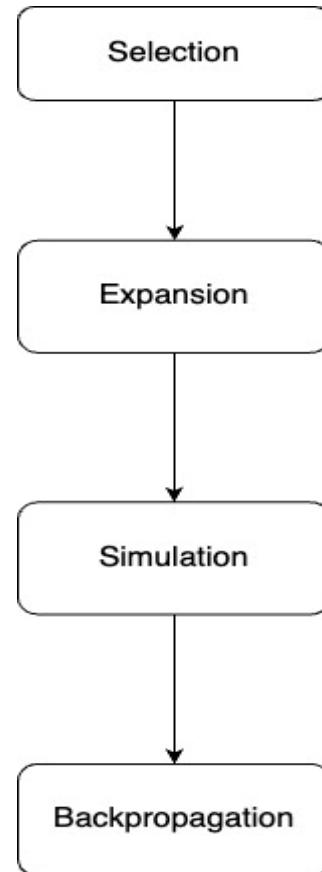
Stage 2 - Expansion Adds new child nodes at the leaf node to represent unexplored actions.

Stage 3 - Simulation: Conducts simulations from the new nodes to estimate outcomes using a default random policy.

Stage 4 - Backpropagation: Updates node values along the selected path based on the results from the simulations to refine future decisions.

Exploration vs. Exploitation

- Balancing Mechanism: The UCB formula during the selection phase is essential for optimizing the balance.
- Upper Confidence Bound (UCB): UCB uses the node's average value and visit count to adjust the search dynamically.



HOW TO ADDRESS THE ISSUE

Addressing Navigational Challenges with MCTS-DRL Integration :

- **Purpose:** The combination of MCTS and DRL is aimed at addressing the complex decision-making required in dynamic environments where both predictability and adaptability are crucial.
- **Implementation:** MCTS structures the decision process through a tree-based exploration model, while DRL provides the learning mechanism that evaluates and optimizes these decisions based on their outcomes.

How MCTS-DRL Addresses the Issue :

- Data Integration: MCTS-DRL processes extensive sensory data from environments to enhance decision-making accuracy.
- Scenario Simulation: Simulates various potential scenarios, enabling proactive planning and obstacle avoidance.
- Informed Decisions: Combines real-time data with learned experiences to make informed decisions that optimize safety and efficiency.
- Dynamic Adaptability: Adjusts strategies in real-time, responding effectively to new obstacles and changing conditions.

DETAILED PROCESS

❖ Data-Driven Decision Making:

- **MCTS:** Explores various potential future scenarios, allowing the robot to consider multiple options.
- **DRL:** Analyzes these options based on historical performance to prioritize the most effective strategies.

❖ Real-Time Adaptation:

- **Feedback Loop:** Uses ongoing operational and environmental data to refine the DRL model, enhancing decision criteria and simulations.
- **Adaptive Learning:** The DRL model adjusts its strategies in response to changing conditions, ensuring effective navigation.

❖ Handling Uncertainties:

- **Probabilistic Exploration:** MCTS probabilistically tests uncertain futures, offering diverse learning opportunities for DRL.
- **Reinforcement Feedback:** DRL refines navigational policies based on outcomes, promoting safe and effective actions.

❖ Enhancing Computational Efficiency:

- **Optimized Exploration:** MCTS focuses on promising paths, optimizing computational efforts.
- **Streamlined Decision-Making:** DRL quickly pinpoints the best actions, minimizing the need for extensive simulations.

WHY MCTS-DRL ENHANCES ROBOTIC NAVIGATION

- **Robust Decision-Making:** MCTS simulates multiple scenarios, DRL optimizes outcomes based on these simulations.
- **Continuous Learning:** System learns and improves from each interaction, enhancing decision accuracy over time.
- **Dynamic Adaptability:** Capable of real-time adjustments to new challenges and environmental changes.
- **Efficient Use of Computational Resources:** MCTS identifies promising decisions, DRL evaluates them quickly, conserving resources.
- **Scalability and Flexibility:** Adapts from simple to complex environments, suitable for various applications like industrial robots and autonomous vehicles.
- **Proven Effectiveness:** Shows marked improvements in navigation efficiency and adaptability in diverse conditions.

EXAMPLE APPLICATION

Case Study: Application of MCTS-DRL in Tesla's Optimus Robot

Tesla's Optimus robot :

Humanoid Design: Optimus is built for dynamic real-world tasks, tailored to handle diverse environments.

MCTS and DRL Integration: Combines Monte Carlo Tree Search with Deep Reinforcement Learning for enhanced navigational and interactive capabilities.

Scenario Simulation: Capable of simulating multiple potential future interaction scenarios in real-time to optimize decision-making.

Adaptive Decision-Making: Enhances the robot's ability to adjust strategies on the fly, crucial for handling new and unexpected situations.

Safety and Efficiency: Ensures high standards of safety while maintaining operational efficiency, crucial for effective task performance.

Real-Time Applications: Designed to interact safely with humans and objects, adapting to complex environments seamlessly.



IMPLEMENTATION

```
class node():
    def __init__(self):
        rospy.init_node('follow', anonymous=False)
        rospy.Subscriber("person_pose_pred_all", PoseArray, self.human_pose_callback, buff_size=1)
        rospy.Subscriber("odom", Odometry, self.odom_callback, buff_size=1)

        self.pub_goal = rospy.Publisher('/move_base_simple/goal', PoseStamped, queue_size =1)
        self.pub_human_pose = rospy.Publisher('/human_pose', Marker, queue_size = 1)

        self.robot_actions, _ = human_traj_t.trajectories(15)
        self.state = np.zeros((2,3))
        self.MCTS_params = {}
        self.MCTS_params['robot_actions'] = self.robot_actions
        file_name = '[seven_actions] [DQN]_1000-layers_1-lsize_16-bs_512-lr_0.01-ep_15-test_1-tar_up100-aSize1-aStep30-2vel'
        model_directory = '/home/sahar/catkin_ws/src/follow-ahead/model/' + file_name + '.pt'
        model = torch.load(model_directory)
        self.MCTS_params['model'] = model
        self.MCTS_params['use_model'] = True
        self.MCTS_params['num_expansion'] = 30
        self.MCTS_params['num_rollout'] = 5
        self.time = time.time()
        self.stay_bool = True

    def odom_callback(self,data):
        robot_p = data.pose.pose.position
        robot_o = data.pose.pose.orientation
        yaw = R.from_quat([0, 0, robot_o.z, robot_o.w]).as_euler('xyz', degrees=False)[2]
        self.state[0,:] = [robot_p.x, robot_p.y, yaw]

    def human_pose_callback(self, data):
        # self.state[0,:] = [0,0,0] #comment this late
        human_traj = np.zeros((6,3))

        for i in range(len(data.poses)-1):
            human_traj[i,:] = [data.poses[i].position.x,
                               data.poses[i].position.y,
                               np.arctan2(data.poses[i+1].position.y-data.poses[i].position.y , data.poses[i+1].position.x-data.poses[i].position.x)]

        human_traj[-1,:] = [data.poses[-1].position.x,
                           data.poses[-1].position.y,
                           human_traj[-2,2]]
```

```
def expand_tree(self, human_traj):
    print()
    print()
    print('current state:')
    print(self.state)

    if not self.stay():
        self.MCTS_params['human_traj'] = human_traj
        nav_state = navState(state=self.state, next_to_move= +1, l = 1)
        node_human = MCTSNode(state=nav_state, params = self.MCTS_params, parent= None, agent = None)
        mcts = MCTS(node_human)
        t1 = time.time()
        robot_best_node, leaf_node = mcts.best_action(0, 0)
        print("expansion time: ", time.time()-t1)
        # robot_best_node, reward, idx = mcts.best_action(0, 0)
        # print("leaf node: ", leaf_node.state.state)
        # self.generate_goal(leaf_node.state.state)
        print('robot"action: ',self.robot_actions[robot_best_node[2]])
        self.generate_goal(robot_best_node[2])
        self.time += 1
    else:
        print("Waiting ...")

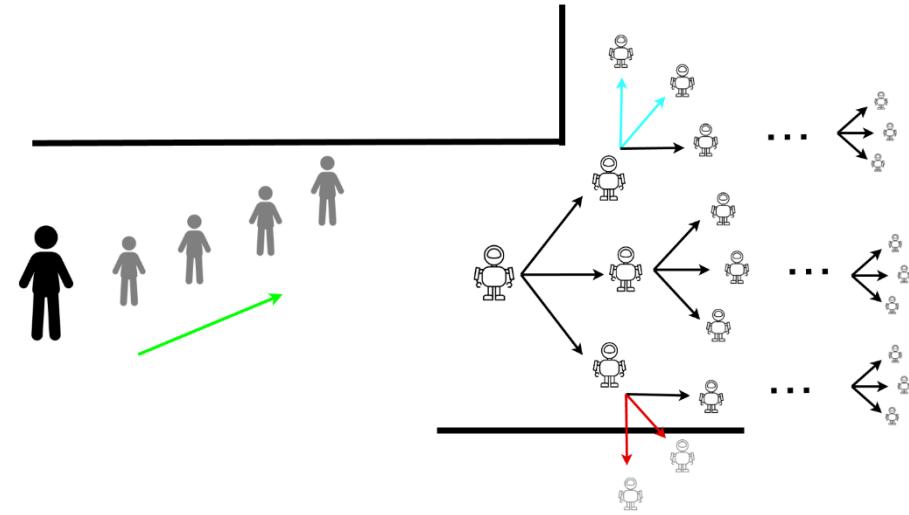
def generate_goal(self, idx):
    TS = 0.4
    theta = self.robot_actions[idx][0] * TS
    D = 3 if self.robot_actions[idx][1] == 0.25 else 3

    goal_robot_frame = [D* np.cos(theta), D* np.sin(theta)]
    print('goal_robot_frame', goal_robot_frame)

    rot_yaw = self.state[0,2]
    trans = self.state[0,:2]
    R_tran = [[np.cos(rot_yaw), -1*np.sin(rot_yaw)], [np.sin(rot_yaw), np.cos(rot_yaw)]]

    goal_map_frame= list(np.add(np.dot(R_tran , goal_robot_frame), trans))
    # goal_map_frame = [goal_state[0,0], goal_state[0,1]]
    print('goal_map_frame', goal_map_frame)
```

OUTPUT



CONCLUSION, LIMITATION, AND FUTURE WORK

❖ Conclusion:

- **Enhanced Navigation:** Demonstrated the efficacy of MCTS and DRL integration in improving robotic navigation in complex environments.
- **Robust Performance:** Proved that these advanced algorithms can adapt to dynamic changes and optimize decision-making.

❖ Future Work:

- **Scaling Up:** Plans to expand the model to multi-agent systems, enhancing coordination and interaction among multiple robots.
- **Enhanced Prediction:** Focus on refining algorithms to improve the prediction accuracy of human intentions for better human-robot interaction.

❖ Limitations:

- **Sensory Dependence:** Relies heavily on high-quality inputs from advanced sensors, impacting performance in sensor-limited scenarios.
- **Computational Intensity:** Faces challenges with the computational demands of processing extensive data in real time.

References

Papers cited	How is the paper used (specifics)	Remarks (if any)
An MCTS-DRL Based Obstacle and Occlusion Avoidance Methodology in Robotic Follow-Ahead Applications	Provided foundational methodology for MCTS and DRL integration in obstacle avoidance.	
A Self-Learning Monte Carlo Tree Search Algorithm for Robot Path Planning	Influenced the development of self-learning algorithms within our robotic path planning module.	
Use of MCTS and DRL for Improving Autonomous Navigation Capabilities in Robots	Supported theoretical underpinnings and practical applications of MCTS and DRL in autonomous navigation.	